

Data Transformation DSL

Problem

- Big data transformation
- Possibly complicated calculations and algorithms

Requirements for a Solution

- scalable
- maintainable
- understandable by non technical person

Approach

- Avoid random access
- Process everything sequentially
- “Sequentially functional”

Why DSL?

- Simple usable interface
- Pluggable implementations
- Declarative approach

Data

- Database is defined as one or more collections
- Collection is an array of Tuples (Rows)
- Tuple has one or more untyped fields

How does it work?

- Chain/Graph of operations
- Similar to Unix pipe just not only linear
- Expression evaluation
- Potential for high degree of parallelization

Basic operations

- define_collection
- generate
- project
- filter
- aggregate
- compose
- group
- sort

`define_collection` collection_name, *fields

Loads a collection (e.g. from a csv file)

generate new_collection, field, count, &block

Generates a new collection with count rows containing one field. The block evaluation determines the value in each row.

project new_collection, collection, options

Can add or remove fields from a collection.
options is map specifying which fields should be included or excluded. It can also contain lambdas to calculate new field values.

shortcut:

calculate new_collection, collection, field, &block

Creates a new collection with just one field per row and value calculated in block.

filter new_collection, collection, &block

Retains rows for which block evaluates to true

aggregate new_collection, collection,
initial_value, field, &block

Performs an aggregation over all rows of a collection returning a new collection with single row and single field

compose new_collection, *collections, &block

It's the classical join, but default is not cartesian product.

group new_collection, collection, options

Groups records according to specification in options.

options[:fields] - fields that drive the group by

options[:computations] - mapping between
new fields and lambdas to calculate those
fields

Sample Implementation

- Calculate impact of discount campaigns
- Leverages exponential smoothing

What is exponential Smoothing?

- popular schema to produce a smoothed Time series
- Single Moving Average - observations weighted equally
- Exponential Smoothing - older observations get exponentially decreasing weights.

Exponential Smoothing

$$s_1 = x_0$$

$$s_t = \alpha x_{t-1} + (1 - \alpha)s_{t-1}, t > 1$$

α is the *smoothing factor*, and $0 \leq \alpha \leq 1$.

The time series look like geometric progression


$$\begin{aligned} s_t &= \alpha x_{t-1} + (1 - \alpha)s_{t-1} \\ &= \alpha x_{t-1} + \alpha(1 - \alpha)x_{t-2} + (1 - \alpha)^2 s_{t-2} \\ &= \alpha [x_{t-1} + (1 - \alpha)x_{t-2} + (1 - \alpha)^2 x_{t-3} + (1 - \alpha)^3 x_{t-4} + \cdots] + (1 - \alpha)^{t-1} x_0. \end{aligned}$$

Optimal Alpha

- sum of the quantities $(s_{n-1} - x_{n-1})^2$ is minimized
- optimal alpha is only valid for a particular sequence
- no correlation to entire data variance


Example

2,1,1,1,1,1,1,1,1,1,1



Example

1,2,1,1,1,1,1,1,1,1,1



Additional problem

- Linear stretching

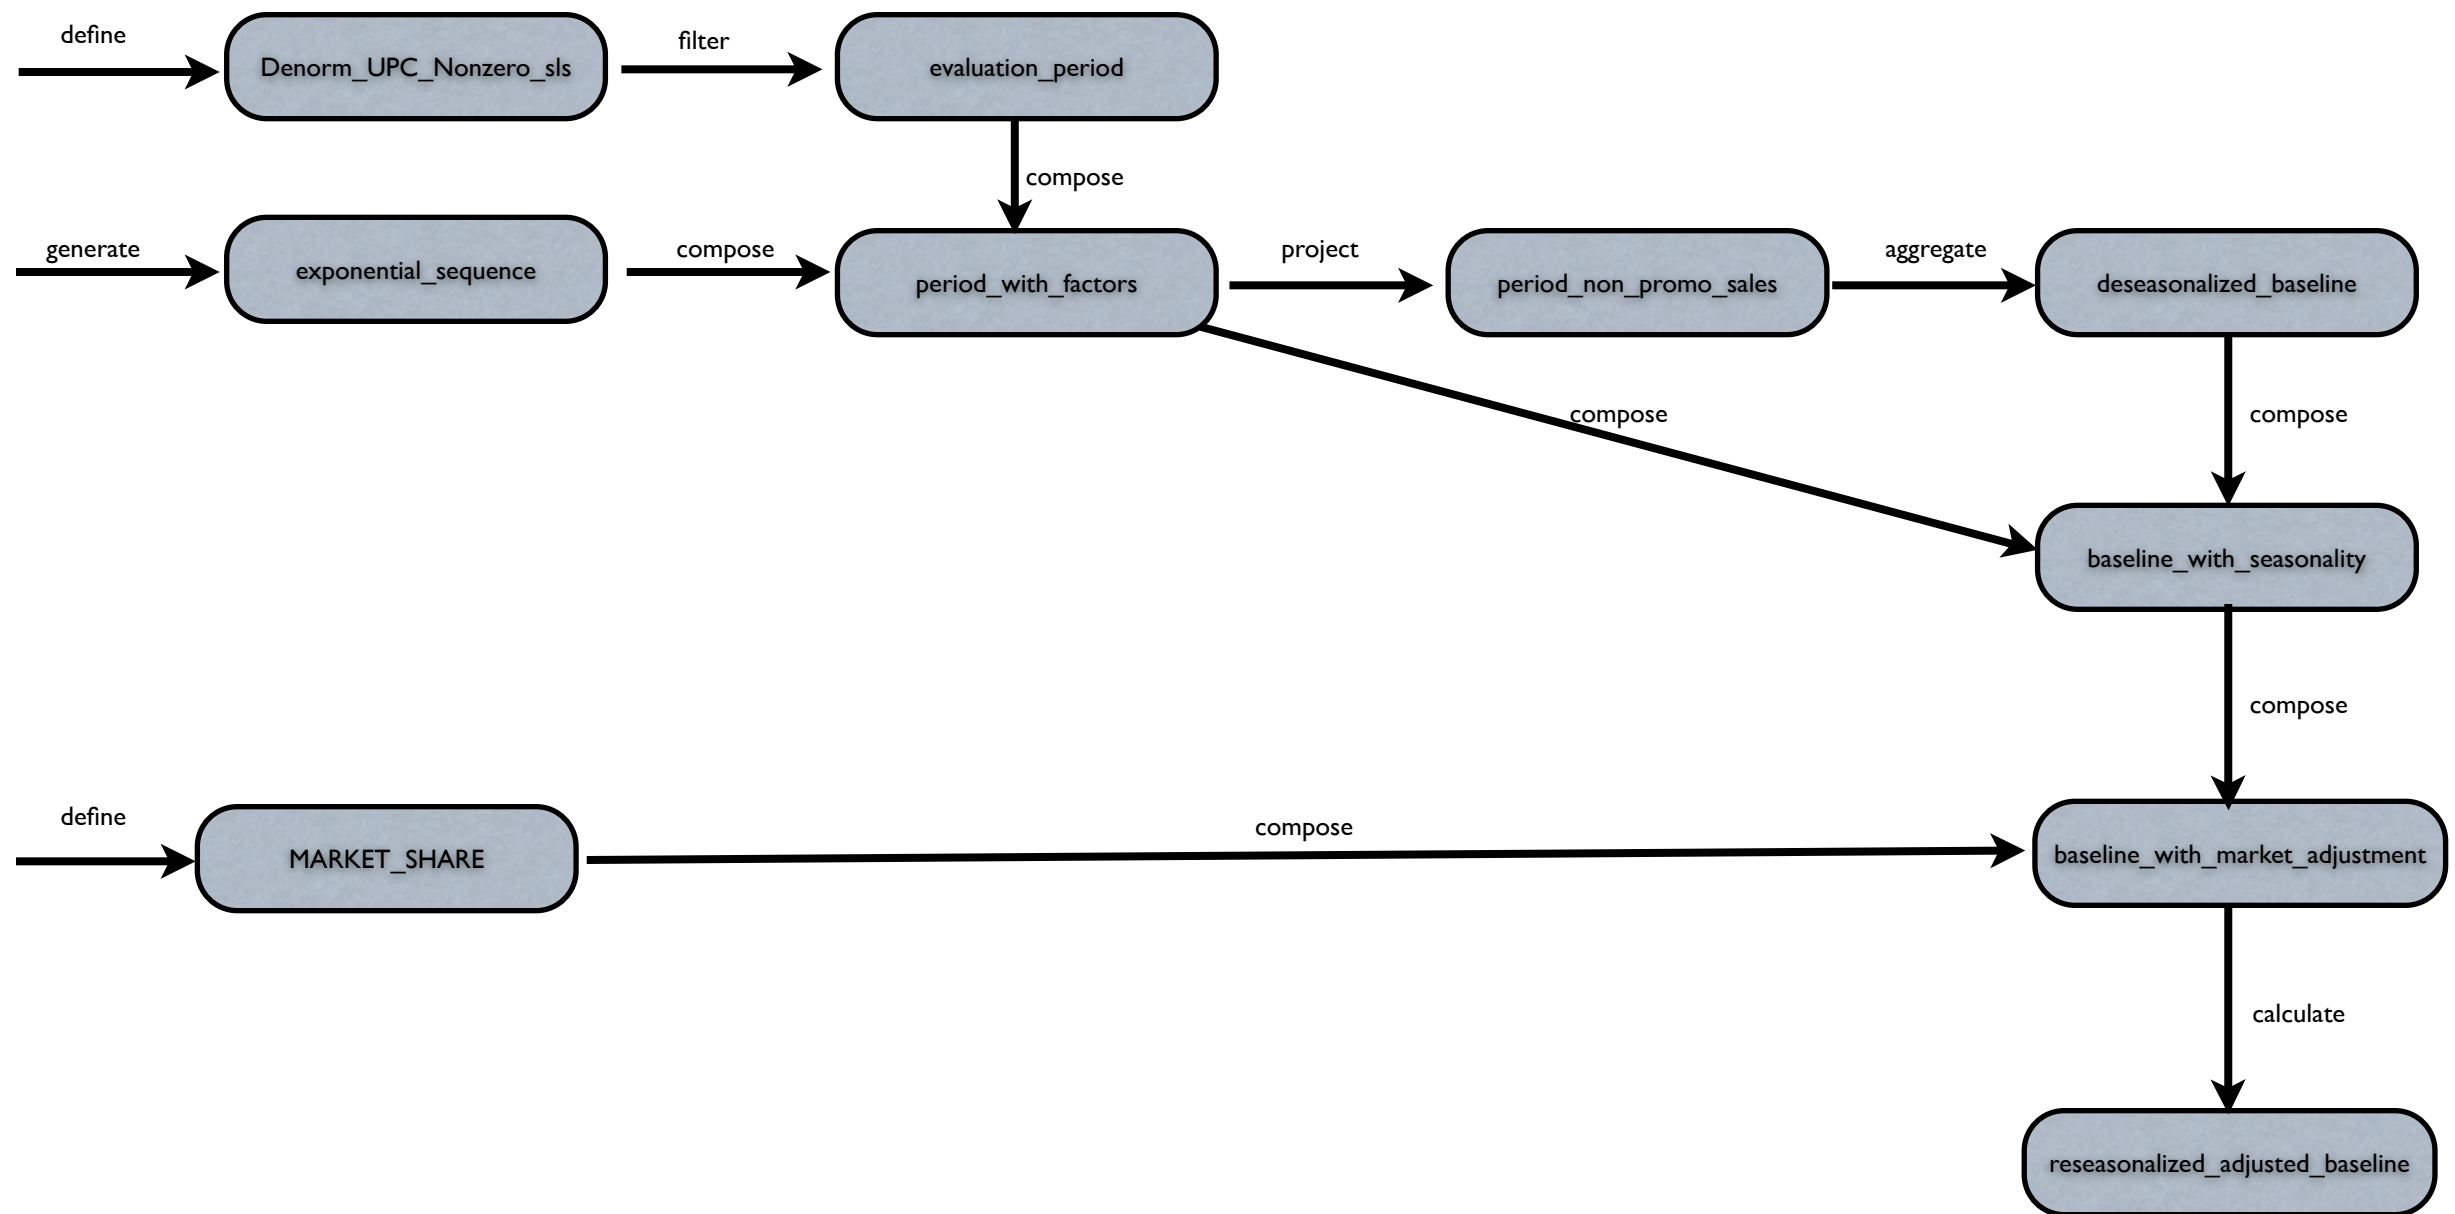
Implementation

- Seemed at first impossible without random access
- Multiple iterations
- How do you do loops in declarative language?
- etc.

The code

```
Transform::Dsl.draw do
  define_collection :Denorm_UPC_Nonzero_sls,
    :UPC_NBR, :WEEK_ID, :FS_SUB_CATEGORY_ID, :FS_SALES_DLRS, :FS_COSTS_DLRS, :FS_PROMO_SALES_DLRS,
    :FS_VENDOR_FUNDING_DLRS, :FS_SALES_UNITS, :FS_PROMO_SALES_UNITS, :FS_NBR_TRANSACTIONS,
    :FS_AVG_BASKET_SIZE, :FS_UNADJUSTED_MARGIN_DLRS, :FS_Lifecycle_stage, :DP_PROMO_FLAG,
    :DC_WEEK_NUM, :DC_YEAR_NUM, :DC_WEEK_DATE, :DC_AD_WEEK, :DS_SUB_CATEGORY_ID, :DS_SEASONALITY_INDEX,
    :FN_start_sale_week, :FN_end_sale_week, :ADJ_SALES_DLRS
  define_collection :MARKET_SHARE, :FS_SUB_CATEGORY_ID, :ADJUSTMENT_FACTOR
  filter(:evaluation_period, :Denorm_UPC_Nonzero_sls) do |sale|
    sale.UPC_NBR.to_i == 28 && (sale.DC_WEEK_NUM.to_i-20).abs <= 8
  end
  generate(:exponential_sequence, :factor, 17) { |i| i==8 ? 0 : (0.5**((i - 8).abs+1))/(1-0.5**8) }
  compose(:period_with_factors, :evaluation_period, :exponential_sequence)
  project(:period_non_promo_sales, :period_with_factors,
    pwk_sales: lambda { |week| week.DP_PROMO_FLAG.to_i == 1 ? week.previous.pwk_sales : week.FS_SALES_DLRS })
  aggregate(:deseasonalized_baseline, :period_non_promo_sales, 0, :deseasonalized_baseline) do |total, week|
    total + week.pwk_sales.to_f*week.factor.to_f/week.DS_SEASONALITY_INDEX.to_f
  end
  compose(:baseline_with_seasonality, :deseasonalized_baseline, :period_with_factors) do |baseline, week|
    week.factor == "0"
  end
  compose(:baseline_with_market_adjustment, :baseline_with_seasonality, :MARKET_SHARE) do |baseline, market|
    baseline.FS_SUB_CATEGORY_ID == market.FS_SUB_CATEGORY_ID
  end
  calculate(:reseasonalized_adjusted_baseline, :baseline_with_market_adjustment, :reseasonalized_adjusted_baseline) do |week|
    week.deseasonalized_baseline.to_f * week.DS_SEASONALITY_INDEX.to_f * week.ADJUSTMENT_FACTOR.to_f
  end
  store :reseasonalized_adjusted_baseline
end
```

Graphical representation



```
define_collection :Denorm_UPC_Nonzero_slr,  
  :UPC_NBR, :WEEK_ID, :FS_SUB_CATEGORY_ID, :FS_SALES_DLRS, :FS_COSTS_DLRS, :FS_PROMO_SALES_DLRS,  
  :FS_VENDOR_FUNDING_DLRS, :FS_SALES_UNITS, :FS_PROMO_SALES_UNITS, :FS_NBR_TRANSACTIONS,  
  :FS_AVG_BASKET_SIZE, :FS_UNADJUSTED_MARGIN_DLRS, :FS_Lifecycle_stage, :DP_PROMO_FLAG,  
  :DC_WEEK_NUM, :DC_YEAR_NUM, :DC_WEEK_DATE, :DC_AD_WEEK, :DS_SUB_CATEGORY_ID, :DS_SEASONALITY_INDEX,  
  :FN_start_sale_week, :FN_end_sale_week, :ADJ_SALES_DLRS
```

```
filter(:evaluation_period, :Denorm_UPC_Nonzero_slts) do |sale|  
  | sale.UPC_NBR.to_i == 28 && (sale.DC_WEEK_NUM.to_i-20).abs <= 8  
end
```

```
generate(:exponential_sequence, :factor, 17) { |i| i==8 ? 0 : (0.5**((i - 8).abs+1))/(1-0.5**8) }
```

```
project(:period_non_promo_sales, :period_with_factors,  
      pwk_sales: lambda { |week| week.DP_PROMO_FLAG.to_i == 1 ? week.previous.pwk_sales : week.FS_SALES_DLRS })
```

```
aggregate(:deseasonalized_baseline, :period_non_promo_sales, 0, :deseasonalized_baseline) do |total, week|  
  total + week.pwk_sales.to_f*week.factor.to_f/week.DS_SEASONALITY_INDEX.to_f  
end
```

```
compose(:baseline_with_seasonality, :deseasonalized_baseline, :period_with_factors) do |baseline, week|  
  | week.factor == "0"  
end
```

```
calculate(:reseasonalized_adjusted_baseline, :baseline_with_market_adjustment, :reseasonalized_adjusted_baseline) do |week|  
  week.deseasonalized_baseline.to_f * week.DS_SEASONALITY_INDEX.to_f * week.ADJUSTMENT_FACTOR.to_f  
end
```



```
store :reseasonalized_adjusted_baseline
```

Demo

Problems

- sort
- compose (join)
- group

Goal

- computational complexity $O(\text{data size})$
- current solution: 10,000 rows/sec

**We are not the only
one**

MongoDB Aggregation Framework

New framework currently in 2.2.0-rc1

\$project

\$match

\$limit

\$skip

\$unwind

\$group

\$sort

**Anyone missing an
operation?**

Anyone missing an
operation?

join

Example Script

```
db.article.aggregate
  { $project : {
    author : 1,
    tags : 1,
  } },
  { $unwind : "$tags" },
  { $group : {
    _id : { tags : 1 },
    authors : { $addToSet : "$author" }
  } }
);
```

Differences

- Only sequential pipe, no graph
- requires de-normalization of data into MongoDB

Questions

- Will users want to use this DSL?
- Is it easy to express problems with it?
- Will we be able to express all cases?
- Can it scale without user paying attention to how the problem is specified?

References

- <http://github.intranet.mckinsey.com/Heinrich-Klobuczek/transform>
- <http://docs.mongodb.org/manual/applications/aggregation/>